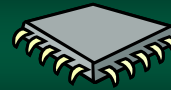




Processors

Part 2

1



Processors

What are they? Besides awesome!

2

Computer Processors

- The *Central Processing Unit (CPU)* is the most complex part of a computer
- In fact, it is the computer!
- It works far different from a high-level language
- *Thousands* of processors have been developed



Spring 2022


Seaworth, Stein - Cook - CS61B

3

3

Some Famous Computer Processors

- RCA 1802
- Intel 8086
- Zilog Z80
- MOS 6502
- Motorola 68000
- ARM



Spring 2022

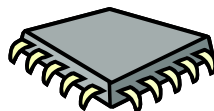
Seaworth, Stein - Cook - CS61B

4

4

Computer Processors

- Each processor functions differently
- Each is designed for a specific purpose – *form follows function*



Spring 2022

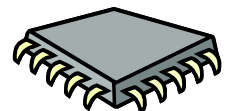
Seaworth, Stein - Cook - CS61B

5

5

Computer Processors

- But all share some basic properties and building blocks...
- Computer hardware is divided into two "units"
 1. Control Logic Unit
 2. Execution Unit



Spring 2022

Seaworth, Stein - Cook - CS61B

6

6

Control Logic Unit (CLU)

- *Control Logic Unit (CLU)* controls the processor
- Determines when instructions can be executed
- Controls internal operations
 - fetch & decode instructions
 - invisible to running programs



Spring 2022

Background: Basic - Cook - CSU 35

7

7

Execution Unit

- *Execution Unit (EU)* contains the hardware that **executes** tasks (your programs)
- Different in many processors
- Modern processors often use multiple execution units to execute instructions in parallel to improve performance

Spring 2022

Background: Basic - Cook - CSU 35

8

8

Execution Unit – The ALU

- *Arithmetic Logic Unit* is part of the Execution Unit and performs all calculations and comparisons
- Processor often contains special hardware for integer and floating point

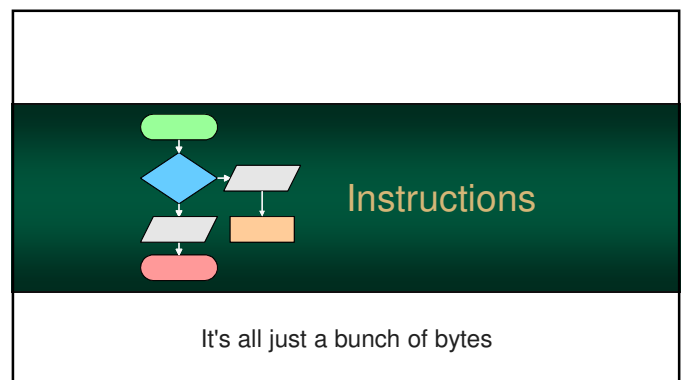


Spring 2022

Background: Basic - Cook - CSU 35

9


9



10

Instructions

- You are used to writing programs in high level programming languages
- Examples:
 - C#
 - Java
 - Python
 - Visual Basic



Spring 2022


Background: Basic - Cook - CSU 35

11

11

High-Level Programming

- These are *third-generation languages*
- They are designed to **isolate** you from architecture of the machine
- This layer of abstraction makes programs "portable" between systems



Spring 2022


Background: Basic - Cook - CSU 35

12

12

Instructions

- Processors **do not** have the constructs you find in high-level languages
- Examples:
 - Blocks
 - If Statements
 - While Statements
 - ... etc



Spring 2022


Secramento State - CS&E - CS&E 35

13

13

Instructions

- Processors can only perform a series of simple tasks
- These are called *instructions*
- Examples:
 - add two values together
 - copy a value
 - jump to a memory location



Spring 2022


Secramento State - CS&E - CS&E 35

14

14

Instructions

- These instructions are used to create all logic needed by a program
- We will cover how to do this during the semester



Spring 2022


Secramento State - CS&E - CS&E 35

15

15

Processor Instruction Set

- A processor's *instruction set* defines all the available instructions
- The instructions and their respective formats are very different for each processor



Spring 2022

Secramento State - CS&E - CS&E 35

16

16

Registers

Where the work is done



Registers

- In high level languages, you put active data into variables
- However, it works quite different on processors
- All computations are performed using *registers*



Spring 2022

Secramento State - CS&E - CS&E 35

18

18

What – exactly – is a register?

- A *register* is a location, on the processor itself, that is used to store temporary data
- Think of it as a special global "variable"
- Some are accessible and usable by a programs, but **many are hidden**



Spring 2022

Backwards State - Cook - CS63B

19

19

What are registers used for?

- Registers are used to store anything the processor needs to keep to track of
- Designed to be *fast!*
- Examples:
 - the result of calculations
 - status information
 - memory location of the running program
 - and much more...

Spring 2022

Backwards State - Cook - CS63B

20

20

General Purpose Registers

- *General Purpose Registers (GPR)* don't have a specific purpose
- They are designed to be used by programs – however they are needed
- Often, you must use registers to perform calculations

Spring 2022

Backwards State - Cook - CS63B

21

21

Special Registers

- There are a number of registers that are used by the Control Logic Unit and cannot be accessed by your program
- This includes registers that control how memory works, your program execution thread, and much more.

Spring 2022

Backwards State - Cook - CS63B

22

22

Special Registers

- *Instruction Pointer (IP)*
 - also called *the program counter*
 - keeps track of the address of your running program
 - think it as the "line number" in your Java program – the one is being executed
 - it can be changed, but only indirectly (*using control logic – which we will cover later*)

Spring 2022

Backwards State - Cook - CS63B

23

23

Special Registers

- *Status Register*
 - contains Boolean information about the processors current state
 - we will use this later, indirectly
- *Instruction Register (IR)*
 - stores the current instruction (being executed)
 - used internally and invisible to your program


Spring 2022

Backwards State - Cook - CS63B

24

24

Register Files



- All the related registers are grouped into a *register file*
- Different processors access and use their register files in very different ways
- Sometimes registers are implied or hardwired

Spring 2022


Secrets: State - Cook - CSU 35

25

25

Components of a Processor

Processor




Spring 2022

Secrets: State - Cook - CSU 35

26

26



Machine Language


The raw bytes of your program

27

27

Machine Language

- The instructions, that are *actually* executed on the processor, are just bytes
- In this raw binary form, instructions are stored in *Machine Language (aka Machine Code)*



Spring 2022


Secrets: State - Cook - CSU 35

28

28

Machine Language

- Each instruction is *encoded* (stored) in a compact binary form
- Easy for the processor to interpret and execute
- Some instructions may take more bytes than others – not all are equal in complexity



Spring 2022

Secrets: State - Cook - CSU 35

29

29

Instruction Encoding

- Each instruction must contain *everything* the processor needs to know to do something
- Think of them as functions in Java: they need a name and arguments to work



Spring 2022

Secrets: State - Cook - CSU 35

30

30

Instruction Encoding

- For example: if you want it to add 2 things...
- The instruction needs:
 - something to tell the processor to add
 - something to identify the two "things"
 - destination to save the result



Spring 2022


Secrecy: State - Cook - CSU 35

31

31

Operation Codes

- Each instruction has a unique operation code (Opcode)
- This is a value that specifies the exact operation to be performed by the processor
- Assemblers use friendly names called *mnemonics*



Spring 2022


Secrecy: State - Cook - CSU 35

32

32

Typical Instruction Format

- The opcode is, typically, followed by various *operands* – what data is to be used
- These can be register codes, addressing data, literal values, etc...




Spring 2022

Secrecy: State - Cook - CSU 35

33

33

Machine Code Example (not x86)



Spring 2022


Secrecy: State - Cook - CSU 35

34

34

Accessing Data

- Processors use registers to hold data being computed
- So, how is data put into the registers to begin with?
- Data can come from two major sources



Spring 2022

Secrecy: State - Cook - CSU 35

36

36

Filling the registers

35

Immediates

- In programming, it is common to assign a constant to a variable
- As you can imagine, this will also be quite common with instructions



Spring 2022

Seacrest, Stein - Cook - CSU 35

37

37

Immediates

- When a constant is stored as part of instruction, it is called an *immediate*
- Once the instruction is loaded by the processor, it is "immediately" available from the IR – hence, the name



Spring 2022

Seacrest, Stein - Cook - CSU 35

38

38

Load Immediate

- A *Load Immediate* instruction, stores a constant into a register
- The instruction must store the destination register and the immediate value



Spring 2022


Seacrest, Stein - Cook - CSU 35

39

39

Load Immediate Example (not x86)

LDI r2, 123




Spring 2022

Seacrest, Stein - Cook - CSU 35

40

40

Copying Data



- Processors have a number of instructions that can copy data
- Each has a unique name –*not surprising since each does something different*

Spring 2022

Seacrest, Stein - Cook - CSU 35

41

41

Transfer

- A *Transfer* instruction, copies the contents of one instruction into another
- The instruction must store both the destination and source register




Spring 2022

Seacrest, Stein - Cook - CSU 35

42

42

Transfer Example (not x86)



43

Effective Addresses

- Processors also have the ability to access memory
- Since memory is a massive array, the processor needs to know the address



44

Effective Addresses


- The *effective address* is used to access memory
- Often it is created by combining multiple values
- ... but we will cover that later in the semester



45

Load


- A *Load* instruction, reads data from memory (at a specified address)
- This data is then stored into the destination register



46


Store

- A *Store* instruction, writes data from a register into the specified address
- Note: the structure is identical to Load



47

Load Example (not x86)




48



Intro to the Intel x64

Part 3

49




The Intel x64

It was simple at first...

50

The Intel x64

- The Intel x64 is the main processor used by servers, laptops, and desktops
- It has evolved continuously over a 40+ year period



Spring 2022


Securix/Securix - Cook - CSO 20

51

51

The Original x86

- First "x86" was the 8086
- Released in 1978
- Attributes:
 - 16-bit registers
 - 16 registers
 - could access of 1MB of RAM (in 64KB blocks using a special "segment" register)



Spring 2022


Securix/Securix - Cook - CSO 20

52

52

What to call the processor

- The classic term "x86" refers to the 32-bit and 16-bit processor family
- With move to 64-bit, the term "x64" is used to differentiate the newest design from the previous



Spring 2022

Securix/Securix - Cook - CSO 20

53

53



Original x86 Registers

It was simple at first...

54

Original x86 Registers

- The original x86 contained 16 registers
- 8 can be used by your programs
- The other 8 are used for memory management



Spring 2022

Background: State - Cook - CSU 35

55

55

Original x86 Registers

- The x86 processor has evolved continuously over the last 4 decades
- It first jumped to 32-bit, and then, again, to 64-bit
- The result is many of the registers have strange names

Spring 2022

Background: State - Cook - CSU 35

56

56

Original x86 Registers

- 8 Registers can be used by your programs
 - Four General Purpose: **AX, BX, CX, DX**
 - Four pointer index: **SI, DI, BP, SP**
- The remaining 8 are restricted
 - Six segment: CS, DS, ES, FS, GS, SS
 - One instruction pointer: IP
 - One status register – used in computations

Spring 2022


Background: State - Cook - CSU 35

57

57

Original General-Purpose Registers

- However, back then (and now too) it is very useful to store 8-bit values
- So, Intel chopped 4 of the registers in half
- These registers have generic names of A, B, C, D



Spring 2022


Background: State - Cook - CSU 35

58

58

Original General-Purpose Registers

- The first and second byte can be used separately or used together
- Naming convention
 - high byte has the suffix "H"
 - low byte has the suffix "L"
 - for both bytes, the suffix is "X"



Spring 2022


Background: State - Cook - CSU 35

59

59

Original General-Purpose Registers

- This essentially doubled the number of registers
- So, there are:
 - four 16-bit registers or
 - eight 8-bit registers
 - ...and any combination you can think off




Spring 2022

Background: State - Cook - CSU 35

60

60

Last the 4 Registers



- The remaining 4 registers were not cut in half
- Used for storing indexes (for arrays) and pointers
- Their purpose
 - DI – destination index
 - SI – source index
 - BP – base pointer
 - SP – stack pointer


Spring 2022

Microprocessors: State - Cook - CSU 35

61

61

Original 16-Bit Registers



Spring 2022

Microprocessors: State - Cook - CSU 35

62

62

Evolution to 64-Bit Registers



This is going to hurt...

Evolution to 32-bit

- When the x86 moved to 32-bit era, Intel expanded the registers to 32-bit
 - the 16-bit ones still exist
 - they have the prefix "e" for extended
- New instructions were added to use them




Spring 2022

Microprocessors: State - Cook - CSU 35

64

64

Original Registers




Spring 2022

Microprocessors: State - Cook - CSU 35

65

65

Expansion to 32-bit




Spring 2022

Microprocessors: State - Cook - CSU 35

66

66

Original Registers




Spring 2022

Background: State - Cook - CSU 20

67

67

Expansion to 32-bit



Spring 2022

Background: State - Cook - CSU 20

68

68

Evolution to 64-bit

- At this point, Intel had decided to abandon the x86 in lieu of their new Itanium Processor
- The Itanium was a radically different design and was completely incompatible
- Advanced Micro Devices (AMD), to Intel's chagrin, decided to — *once again* — extend the x86



Spring 2022

Background: State - Cook - CSU 20

69

69

Evolution to 64-bit

- Registers were extended again
 - 64-bit registers have the prefix "*r*"
 - 8 additional registers were added
 - also, it is now possible to get 8-bit values from all registers (hardware is more consistent!)
- Some old, archaic, features were dropped



Spring 2022

Background: State - Cook - CSU 20

70

70

Evolution to 64-bit

- The AMD-64 was a huge commercial *success*
- The Itanium was a commercial *failure*
- Intel, dropped the Itanium and started making 64-bit x86 using AMD's design




Spring 2022

Background: State - Cook - CSU 20

71

71

Expansion to 64-bit




Spring 2022

Background: State - Cook - CSU 20

72

72

Expansion to 64-bit




Spring 2022

Section 2.1.1: State - C++ - CSU 38

73

73

Expansion to 64-bit




Spring 2022

Section 2.1.1: State - C++ - CSU 35

74

74

Expansion to 64-bit




Spring 2022

Section 2.1.1: State - C++ - CSU 35

75

75

New 64-bit Registers: R8...R15



Spring 2022

Section 2.1.1: State - C++ - CSU 35

76

76

64-Bit Register Table

Register	32 bit	16-bit	8-bit High	8-bit Low
rax	eax	ax	ah	al
rbx	ebx	bx	bh	bl
rcx	ecx	cx	ch	cl
rdx	edx	dx	dh	dl
rsi	esi	si		sil
rdi	edi	di		dil
rbp	ebp	bp		bpl
rsp	esp	sp		spl

Spring 2022

Section 2.1.1: State - C++ - CSU 35

77

77

64-Bit Register Table

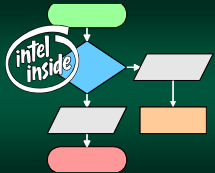
Register	32 bit	16-bit	8-bit High	8-bit Low
r8	r8d	r8w		r8b
r9	r9d	r9w		r9b
r10	r10d	r10w		r10b
r11	r11d	r11w		r11b
r12	r12d	r12w		r12b
r13	r13d	r13w		r13b
r14	r14d	r14w		r14b
r15	r15d	r15w		r15b

Spring 2022

Section 2.1.1: State - C++ - CSU 35

78

78



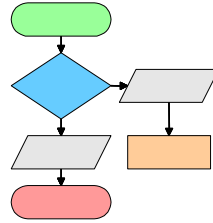
Basic Intel x86 Instructions

Feel the pow-wah of the x86!

79

Basic Intel x86 Instructions

- Each x86 instruction can have up to 2 operands
- Operands in x86 instructions are very versatile
- Each operand can be either a memory address, register or an immediate value



Spring 2022 Backwards State - Cook - CS6 35 80

80

Types of Operands

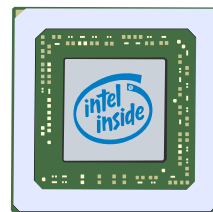
- Registers
- Address in memory
- Register pointing to a memory address
- Immediate

Spring 2022 Backwards State - Cook - CS6 35 81

81

Intel x86 Instruction Limits

- There are some limitations...
- Some instructions must use an immediate
- Some instructions require a specific register to perform calculations



Spring 2022 Backwards State - Cook - CS6 35 82

82

Intel x86 Instruction Limits

- A register must always be involved
 - processors use registers for all activity
 - both operands cannot access memory at the same time
 - the processor has to have it at some point!*
- Also, obviously, the receiving field cannot be an immediate value

Spring 2022 Backwards State - Cook - CS6 35 83

83

Instruction: Move

- The Intel Move Instruction combines transfer, load and store instructions under one name
- ... well, that's something the assembler does for us – but, we'll cover that soon
- "Move" is a tad confusing – it copies data

Spring 2022 Backwards State - Cook - CS6 35 84

84

Instruction: Move

```
MOV destination, source
```

Immediate, Register, Memory

Register, Memory

85

Example: Move immediate

```
MOV rax, 42
```

Source is a immediate constant

Same as Java
`rax = 42;`

Destination is rax

86

Example: Transfer

```
MOV rbx, rax
```

Source is rax

Same as Java
`rbx = rax;`

Destination is rbx

87

Example: Load

```
MOV rax, total
```

"total" is memory location

Destination is rax

88

Example: Store

```
MOV counter, rax
```

Source is rax

Memory location named 'Counter'

89

Example: "A" Register

So many options!

```
mov al, 42      #low byte
mov ah, 13      #high byte
mov ax, 1947    #both bytes
```

90

Instruction: Add & Subtract

- The Add and Subtract instructions take two operands and store the result in the first operand
- This is the same as the **+=** and **-=** operators used in Visual Basic .NET, C, C++, Java, etc...



Spring 2022

Backwards State - Cook - CS6 35

91

91

Instruction: Add

ADD *target* , *value*

Immediate, Register, Memory

Register, Memory

Spring 2022

Backwards State - Cook - CS6 35

92

92

Example: Move register to memory

Move memory into rax

```
MOV rax, counter
ADD rax, 2
```

Same as Java
`rax += 2;`

Spring 2022


Backwards State - Cook - CS6 35

93

93

Instruction: And & Or

- The Bitwise And & Bitwise Or instructions take two operands and stores the result in the second operand
- This is the same as the **&=** and **|=** operators used in C, C++, Java, etc...



Spring 2022

Backwards State - Cook - CS6 35

94

94

Instruction: Logical And

AND *target* , *value*

Immediate, Register, Memory

Register, Memory

Spring 2022

Backwards State - Cook - CS6 35

95

95

Example: Logical Or

```
#Convert 5 to ASCII '5'
MOV rax, 5
OR rax, 0x30
```

0011 0000

Spring 2022

Backwards State - Cook - CS6 35

96

96

Call Instruction

- The *Call Instruction* causes the processor to start running instructions at a specified memory location (a subroutine)
- Subroutines are analogous to the functions you wrote in Java
- Once it completes, execution returns from the subroutine and continues after the call

Spring 2022

Backwards State - CS635

97

97

Call Instruction

CALL *address*

Usually a label – a constant that holds an address

Spring 2022


Backwards State - CS635

98

98

Call Instruction

- The *Call instruction* doesn't change any of the general-purpose registers
- It only stores an address – where execution will continue



Spring 2022

Backwards State - CS635

99

99

Example: Print an integer

#Using the CSC35 library

MOV rcx, 1846

CALL PrintInt

This name is an address

Spring 2022

Backwards State - CS635

100

100