



## Buffers & Direct Storage

### Part 5

1



## Buffers

Creating your own space

2

## Buffers

- A *buffer* is any allocated block of memory that contains data
- This can hold anything:
  - text
  - image
  - file
  - etc....



Fall 2020

Sacramento State - Cook - CSc 35

3

3

## Buffers



- There are several assembly **directives** which will allocate space
- We have covered a few of them, but there are many – all with a specific purpose

Fall 2020

Sacramento State - Cook - CSc 35

4

4

## A few directives that create space

Directive	What it does
<b>.ascii</b>	Allocate enough space to store an ASCII string
<b>.quad</b>	Allocate 8-byte blocks with initial value(s)
<b>.byte</b>	Allocate byte(s) with initial value(s)
<b>.space</b>	Allocate any <b>size</b> of empty bytes (with initial values).

Fall 2020

Sacramento State - Cook - CSc 35

5

5

## Labels are addresses

- Labels are used to keep track of memory locations
- They are stored, by the assembler, in a table
- Whenever a label is used in the program, the assembler substitutes the address



Fall 2020

Sacramento State - Cook - CSc 35

6

6

## Labels are addresses

- The table of labels is stored in the *object file*
- That way the linker can resolve any unknown labels
- After the program is linked into an executable, only addresses exist. No labels.

MY NAME IS



Fall 2020 Sacramento State - Cook - CSc 35 7

7

## Quad Directive

Let's assume Value = 2000

Value:  
.quad 74

2000	4A
2001	00
2002	00
2003	00
2004	00
2005	00
2006	00
2007	00

Fall 2020 Sacramento State - Cook - CSc 35 8

8

## ASCII Directive Creates a Buffer

Text:

.ascii "Hello\0"

This label will store an address... once the assembler finds where to store it.

Creates 6 bytes to store Hello. They are stored consecutively.

Fall 2020 Sacramento State - Cook - CSc 35 9

9

## Bytes are stored consecutively

Let's assume Text = 2000

Text:  
.ascii "Hello\0"

2000	48	H
2001	65	e
2002	6C	l
2003	6C	l
2004	6F	o
2005	00	\0

Fall 2020 Sacramento State - Cook - CSc 35 10

10

## Same Thing!

Text:

.byte 'H'  
.byte 'e'  
.byte 'l'  
.byte 'l'  
.byte 'o'  
.byte '\0'

Created byte by byte

Fall 2020 Sacramento State - Cook - CSc 35 11

11

## This works too!

Text:  
.ascii "Hello"  
.byte 0

Directives just create space. So, this creates a byte after the ASCII text.

Fall 2020 Sacramento State - Cook - CSc 35 12

12

## Create a Buffer of Any Size

Text :

.space 30

Create 30 bytes  
(defaults to 0x20  
which is a space)

Fall 2020

Sacramento State - Cook - CSc 35

13

## Create a Buffer of Any Size

Text :

.space 30, 0

Create 30 bytes.  
All of which are 0.

Fall 2020

Sacramento State - Cook - CSc 35

14



## Direct Addressing

Using memory... finally

15

## Direct Addressing


- In *direct addressing*, the processor reads data directly from the an address
- Commonly used to:
  - get a value from a "variable"
  - read items in an array
  - etc...

Fall 2020

Sacramento State - Cook - CSc 35

16

## Direct Addressing



Fall 2020

Sacramento State - Cook - CSc 35

17

## Direct in Java

- The following, for comparison, is the equivalent code in Java
- The memory at the address *total* is loaded into *rax*

```
// rax = Memory[total];
mov rax, total
```

Fall 2020

Sacramento State - Cook - CSc 35

18

## LEA vs MOV

- Load Effective Address stores an address into a register
- For Direct Addressing, the address is sent to the bus (to access memory)

```
// rax = total;  
lea rax, total
```

Fall 2020 Sacramento State - Cook - CSc 35

19

19

## Example: Direct

```
.intel_syntax noprefix  
.data  
funds:  
    .quad 100  
  
.text  
.global _start  
_start:  
    mov rbx, funds
```

64 bit integer  
with an initial value of 100.

Read 8 bytes at this address.  
Doesn't store "the" address in rbx.

Fall 2020 Sacramento State - Cook - CSc 35

20

20

## Direct in Java

- **Note:** this a shortcut notation
- The full notation would use square brackets
- The assembler recognizes the difference automatically

```
// rax = Memory[total];  
mov rax, total
```

Fall 2020 Sacramento State - Cook - CSc 35

21

21

## Direct in Java

- You can use the square-brackets if you want
- This way it explicitly show *how* the label is being used – it's a matter of preference

```
// rax = Memory[total];  
mov rax, [total]
```

Fall 2020 Sacramento State - Cook - CSc 35

22

22

## Example: Direct

```
.intel_syntax noprefix  
.data  
funds:  
    .quad 100  
  
.text  
.global _start  
_start:  
    mov rax, [funds]
```

A bit more descriptive


Fall 2020 Sacramento State - Cook - CSc 35

23

23

## Cause of the Segmentation Fault

- Knowing when to use an address or the data *located at that address* is vital
- This is one of the most common mistakes is programming



Fall 2020 Sacramento State - Cook - CSc 35

24

24

## Cause of the Segmentation Fault

```
.intel_syntax noprefix
.data
Message:
    .ascii "Hello!!\0"
```

Creates 8 bytes using  
ASCII values

```
.text
.global _start
_start:
    mov    rbx, Message
    call   PrintCString
```

Used mov rather than lea.  
rbx is 64-bit (8 bytes)

Fall 2020 Sacramento State - Cook - CSc 35

25

25

## Cause of the Segmentation Fault

```
Message:
    .ascii "Hello!!\0"

.text
.global _start
_start:
    mov    rbx, Message
    call   PrintCString
```

Message	48	H
	65	e
	6C	l
	6C	l
	6F	o
	21	!
	21	!
	00	\0

Fall 2020 Sacramento State - Cook - CSc 35

26

26

## Cause of the Segmentation Fault

```
Message:
    .ascii "Hello!!\0"

.text
.global _start
_start:
    mov    rbx, Message
    call   PrintCString
```

Grabs 8 bytes and  
creates a huge value

Message	48	H
	65	e
	6C	l
	6C	l
	6F	o
	21	!
	21	!
	00	\0

Fall 2020 Sacramento State - Cook - CSc 35

27

27

## Cause of the Segmentation Fault

```
Message:
    .ascii "Hello!!\0"

.text
.global _start
_start:
    lea    rbx, Message
    call   PrintCString
```

PrintCString needs the  
address of 'Message'

Fall 2020 Sacramento State - Cook - CSc 35

28

28

## Sizing Instructions



How many bytes are you using?

## Sizing Instructions

- The Intel can load/store 1-byte, 2-byte, 4-byte or 8-byte values
- The assembler knows (*by looking at the size of the register*) how much many bytes you want to load/store



Fall 2020 Sacramento State - Cook - CSc 35


29

30

29

## Sizing Instructions

- However, sometimes the number of bytes (1, 2, etc..) can't be determined
- In this case, the assembler will report an error
- ... since it doesn't know how to encode the instruction



Fall 2020

Sacramento State - Cook - CSc 35

31

31

## Example: How Many Bytes?

```
.intel_syntax noprefix
.data
total:
    .quad 0

.text
.global _start
_start:
    mov total, 50
```

total is a target address.  
It doesn't have any implied size.

Fall 2020

Sacramento State - Cook - CSc 35

32

32

## Example: How Many Bytes?

```
.intel_syntax noprefix
.data
total:
    .quad 0

.text
.global _start
_start:
    mov total, 50
```

How many bytes is this?  
The value 50 can be stored  
in 1, 2, 4, or 8 bytes.

Fall 2020

Sacramento State - Cook - CSc 35

33

33

## How Many Bytes?

- If it is not obvious to the assembler how many bytes you want to access, it will report *"ambiguous operand size"*
- To address this issue...
  - GAS assembly allows you place a single character after the instruction's mnemonic
  - this suffix will tell the assembler how many bytes will be accessed during the operation

Fall 2020

Sacramento State - Cook - CSc 35

34

34

## How Many Bytes

Suffix	Name	Size
<b>b</b>	byte	1 byte
<b>s</b>	short	2 bytes
<b>l</b>	long	4 bytes
<b>q</b>	quad	8 bytes

Fall 2020

Sacramento State - Cook - CSc 35

35

35

## Example: Suffix Used

```
.intel_syntax noprefix
.data
total:
    .quad 0

.text
.global _start
_start:
    movq total, 50
```

Now the assembler knows  
you mean "move quad".


Fall 2020

Sacramento State - Cook - CSc 35

36

36

# Endianness




The "proper" order of things

37

# So Many Bytes...

- On a 64-bit system, each word consists of 8 bytes
- So, when any 64-bit value is stored in memory, each of those 8 bytes must be stored
- However, question remains: *What order do we store them?*



38

# Example Unsigned Integer (4 Byte)

1,188,852,977

46	DC	74	F1
----	----	----	----

Most significant Byte (MSB)

Least significant Byte (LSB)

39

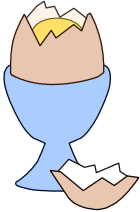
# So Many Bytes...

- Do we store the least-significant byte (LSB) first, or the most-significant (MSB)?
- As long as a system always follows the same format, then there are no problems
- ... but different system use different approaches

40

# Big Endian vs. Little Endian

- Big-Endian approach
  - store the MSB first
  - used by Motorola & PowerPC
- Little-Endian approach
  - store the LSB first
  - used by Intel
  - appears "backwards" in hex-editors



41

# Big Endian vs. Little Endian

46	DC	74	F1
----	----	----	----

**Big Endian**


0	46
1	DC
2	74
3	F1

**Little Endian**

0	F1
1	74
2	DC
3	46

42

## Assuming Value is at 2000



43

## No "End" to Problems

- There is a problem... if two systems use different formats, data will be interpreted incorrectly!
- If how the read differs from how it is stored, the data will be mangled



44

## No "End" to Problems

- For example:
  - a little-endian system reads a value stored in big-endian
  - a big-endian system reads a value stored in little-endian
- Programmers must be conscience of this whenever binary data is accessed



45

## No "End" to Problems

- So, whenever data is read from secondary storage, you cannot assume it will be in your processor's format
- This is compounded by file formats (gif, jpeg, mp3, etc...) which are also inconsistent



46


## Example File Format Endianness

File Format	Endianness
Adobe Photoshop	Big Endian
Windows Bitmap (.bmp)	Little Endian
GIF	Little Endian
JPEG	Big Endian
MP4	Big Endian
ZIP file	Little Endian

47

## So... who is correct?

- So, what is the correct and superior format?
- Is it Intel (little endian)?
- ...or the PowerPC (big endian) correct?




48



## So... who is correct?

- In reality neither side is superior
- Both formats are equally correct
- Both have minor advantages in assembly... but nothing huge



Fall 2020

Sacramento State - Cook - CS& 35

49

49

## Gulliver's Travels



Fall 2020

Sacramento State - Cook - CS& 35

50

50